

Building a Recurrent Neural Network - Step by Step - v3

July 20, 2019

1 Building your Recurrent Neural Network - Step by Step

Welcome to Course 5's first assignment! In this assignment, you will implement your first Recurrent Neural Network in numpy.

Recurrent Neural Networks (RNN) are very effective for Natural Language Processing and other sequence tasks because they have “memory”. They can read inputs $x^{(t)}$ (such as words) one at a time, and remember some information/context through the hidden layer activations that get passed from one time-step to the next. This allows a uni-directional RNN to take information from the past to process later inputs. A bidirection RNN can take context from both the past and the future.

Notation: - Superscript $[l]$ denotes an object associated with the l^{th} layer. - Example: $a^{[4]}$ is the 4th layer activation. $W^{[5]}$ and $b^{[5]}$ are the 5th layer parameters.

- Superscript (i) denotes an object associated with the i^{th} example.
 - Example: $x^{(i)}$ is the i^{th} training example input.
- Superscript $\langle t \rangle$ denotes an object at the t^{th} time-step.
 - Example: $x^{(t)}$ is the input x at the t^{th} time-step. $x^{(i)\langle t \rangle}$ is the input at the t^{th} timestep of example i .
- Lowercase i denotes the i^{th} entry of a vector.
 - Example: $a_i^{[l]}$ denotes the i^{th} entry of the activations in layer l .

We assume that you are already familiar with numpy and/or have completed the previous courses of the specialization. Let's get started!

Let's first import all the packages that you will need during this assignment.

```
In [15]: import numpy as np
         from rnn_utils import *
```

1.1 1 - Forward propagation for the basic Recurrent Neural Network

Later this week, you will generate music using an RNN. The basic RNN that you will implement has the structure below. In this example, $T_x = T_y$.

Figure 1: Basic RNN model

Here's how you can implement an RNN:

Steps: 1. Implement the calculations needed for one time-step of the RNN. 2. Implement a loop over T_x time-steps in order to process all the inputs, one at a time.

Let's go!

1.2 1.1 - RNN cell

A Recurrent neural network can be seen as the repetition of a single cell. You are first going to implement the computations for a single time-step. The following figure describes the operations for a single time-step of an RNN cell.

Figure 2: Basic RNN cell. Takes as input $x^{(t)}$ (current input) and $a^{(t-1)}$ (previous hidden state containing information from the past), and outputs $a^{(t)}$ which is given to the next RNN cell and also used to predict $y^{(t)}$

Exercise: Implement the RNN-cell described in Figure (2).

Instructions: 1. Compute the hidden state with tanh activation: $a^{(t)} = \tanh(W_{aa}a^{(t-1)} + W_{ax}x^{(t)} + b_a)$. 2. Using your new hidden state $a^{(t)}$, compute the prediction $\hat{y}^{(t)} = \text{softmax}(W_{ya}a^{(t)} + b_y)$. We provided you a function: `softmax`. 3. Store $(a^{(t)}, a^{(t-1)}, x^{(t)}, \text{parameters})$ in cache 4. Return $a^{(t)}, y^{(t)}$ and cache

We will vectorize over m examples. Thus, $x^{(t)}$ will have dimension (n_x, m) , and $a^{(t)}$ will have dimension (n_a, m) .

In [16]: # GRADED FUNCTION: `rnn_cell_forward`

```
def rnn_cell_forward(xt, a_prev, parameters):
    """
    Implements a single forward step of the RNN-cell as described in Figure (2)

    Arguments:
    xt -- your input data at timestep "t", numpy array of shape (n_x, m)
    a_prev -- Hidden state at timestep "t-1", numpy array of shape (n_a, m)
    parameters -- python dictionary containing:
                    Wax -- Weight matrix multiplying the input, numpy array of shape (n_a, n_x)
                    Waa -- Weight matrix multiplying the hidden state, numpy array of shape (n_a, n_a)
                    Wya -- Weight matrix relating the hidden-state to the output, numpy array of shape (n_y, n_a)
                    ba -- Bias, numpy array of shape (n_a, 1)
                    by -- Bias relating the hidden-state to the output, numpy array of shape (n_y, 1)

    Returns:
    a_next -- next hidden state, of shape (n_a, m)
    yt_pred -- prediction at timestep "t", numpy array of shape (n_y, m)
    cache -- tuple of values needed for the backward pass, contains (a_next, a_prev, xt)
    """

    # Retrieve parameters from "parameters"
    Wax = parameters["Wax"]
    Waa = parameters["Waa"]
    Wya = parameters["Wya"]
    ba = parameters["ba"]
    by = parameters["by"]

    ### START CODE HERE ### (2 lines)
    # compute next activation state using the formula given above
    a_next = np.tanh( np.dot(Waa, a_prev) + np.dot(Wax, xt) + ba)
    # compute output of the current cell using the formula given above
```

```

yt_pred = softmax(np.dot(Wya, a_next) + by)
### END CODE HERE ###

# store values you need for backward propagation in cache
cache = (a_next, a_prev, xt, parameters)

return a_next, yt_pred, cache

```

```

In [17]: np.random.seed(1)
         xt = np.random.randn(3,10)
         a_prev = np.random.randn(5,10)
         Waa = np.random.randn(5,5)
         Wax = np.random.randn(5,3)
         Wya = np.random.randn(2,5)
         ba = np.random.randn(5,1)
         by = np.random.randn(2,1)
         parameters = {"Waa": Waa, "Wax": Wax, "Wya": Wya, "ba": ba, "by": by}

         a_next, yt_pred, cache = rnn_cell_forward(xt, a_prev, parameters)
         print("a_next[4] = ", a_next[4])
         print("a_next.shape = ", a_next.shape)
         print("yt_pred[1] =", yt_pred[1])
         print("yt_pred.shape = ", yt_pred.shape)

a_next[4] = [ 0.59584544  0.18141802  0.61311866  0.99808218  0.85016201  0.99980978
 -0.18887155  0.99815551  0.6531151   0.82872037]
a_next.shape = (5, 10)
yt_pred[1] = [ 0.9888161   0.01682021  0.21140899  0.36817467  0.98988387  0.88945212
 0.36920224  0.9966312   0.9982559   0.17746526]
yt_pred.shape = (2, 10)

```

Expected Output:

```

a_next[4]:
[ 0.59584544 0.18141802 0.61311866 0.99808218 0.85016201 0.99980978 -0.18887155 0.99815551
0.6531151 0.82872037]
a_next.shape:
(5, 10)
yt[1]:
[ 0.9888161 0.01682021 0.21140899 0.36817467 0.98988387 0.88945212 0.36920224 0.9966312
0.9982559 0.17746526]
yt.shape:
(2, 10)

```

1.3 1.2 - RNN forward pass

You can see an RNN as the repetition of the cell you've just built. If your input sequence of data is carried over 10 time steps, then you will copy the RNN cell 10 times. Each cell takes as input the

hidden state from the previous cell ($a^{(t-1)}$) and the current time-step's input data ($x^{(t)}$). It outputs a hidden state ($a^{(t)}$) and a prediction ($y^{(t)}$) for this time-step.

Figure 3: Basic RNN. The input sequence $x = (x^{(1)}, x^{(2)}, \dots, x^{(T_x)})$ is carried over T_x time steps. The network outputs $y = (y^{(1)}, y^{(2)}, \dots, y^{(T_x)})$.

Exercise: Code the forward propagation of the RNN described in Figure (3).

Instructions: 1. Create a vector of zeros (a) that will store all the hidden states computed by the RNN. 2. Initialize the "next" hidden state as a_0 (initial hidden state). 3. Start looping over each time step, your incremental index is t : - Update the "next" hidden state and the cache by running `rnn_cell_forward` - Store the "next" hidden state in a (t^{th} position) - Store the prediction in y - Add the cache to the list of caches 4. Return a , y and caches

In [18]: # GRADED FUNCTION: `rnn_forward`

```
def rnn_forward(x, a0, parameters):
    """
    Implement the forward propagation of the recurrent neural network described in Figure 3.

    Arguments:
    x -- Input data for every time-step, of shape (n_x, m, T_x).
    a0 -- Initial hidden state, of shape (n_a, m)
    parameters -- python dictionary containing:
                    Waa -- Weight matrix multiplying the hidden state, numpy array of shape (n_a, n_a)
                    Wax -- Weight matrix multiplying the input, numpy array of shape (n_a, m)
                    Wya -- Weight matrix relating the hidden-state to the output, numpy array of shape (n_y, n_a)
                    ba -- Bias numpy array of shape (n_a, 1)
                    by -- Bias relating the hidden-state to the output, numpy array of shape (n_y, 1)

    Returns:
    a -- Hidden states for every time-step, numpy array of shape (n_a, m, T_x)
    y_pred -- Predictions for every time-step, numpy array of shape (n_y, m, T_x)
    caches -- tuple of values needed for the backward pass, contains (list of caches, x)
    """

    # Initialize "caches" which will contain the list of all caches
    caches = []

    # Retrieve dimensions from shapes of x and parameters["Wya"]
    n_x, m, T_x = x.shape
    n_y, n_a = parameters["Wya"].shape

    ### START CODE HERE ###

    # initialize "a" and "y" with zeros (2 lines)
    a = np.zeros(shape=(n_a, m, T_x))
    y_pred = np.zeros(shape=(n_y, m, T_x))

    # Initialize a_next (1 line)
    a_next = a0
```

```

    # loop over all time-steps
    for t in range(T_x):
        # Update next hidden state, compute the prediction, get the cache (1 line)
        a_next, yt_pred, cache = rnn_cell_forward(x[... , t], a_next, parameters)
        # Save the value of the new "next" hidden state in a (1 line)
        a[:, :, t] = a_next
        # Save the value of the prediction in y (1 line)
        y_pred[:, :, t] = yt_pred
        # Append "cache" to "caches" (1 line)
        caches.append(cache)

    ### END CODE HERE ###

    # store values needed for backward propagation in cache
    caches = (caches, x)

    return a, y_pred, caches

```

```

In [19]: np.random.seed(1)
x = np.random.randn(3,10,4)
a0 = np.random.randn(5,10)
Waa = np.random.randn(5,5)
Wax = np.random.randn(5,3)
Wya = np.random.randn(2,5)
ba = np.random.randn(5,1)
by = np.random.randn(2,1)
parameters = {"Waa": Waa, "Wax": Wax, "Wya": Wya, "ba": ba, "by": by}

a, y_pred, caches = rnn_forward(x, a0, parameters)
print("a[4][1] = ", a[4][1])
print("a.shape = ", a.shape)
print("y_pred[1][3] =", y_pred[1][3])
print("y_pred.shape = ", y_pred.shape)
print("caches[1][1][3] =", caches[1][1][3])
print("len(caches) = ", len(caches))

a[4][1] = [-0.99999375  0.77911235 -0.99861469 -0.99833267]
a.shape = (5, 10, 4)
y_pred[1][3] = [ 0.79560373  0.86224861  0.11118257  0.81515947]
y_pred.shape = (2, 10, 4)
caches[1][1][3] = [-1.1425182 -0.34934272 -0.20889423  0.58662319]
len(caches) = 2

```

Expected Output:

```

a[4][1]:
[-0.99999375 0.77911235 -0.99861469 -0.99833267]

```

```

a.shape:
(5, 10, 4)
y[1][3]:
[ 0.79560373 0.86224861 0.11118257 0.81515947]
y.shape:
(2, 10, 4)
cache[1][1][3]:
[-1.1425182 -0.34934272 -0.20889423 0.58662319]
len(cache):
2

```

Congratulations! You’ve successfully built the forward propagation of a recurrent neural network from scratch. This will work well enough for some applications, but it suffers from vanishing gradient problems. So it works best when each output $y^{(t)}$ can be estimated using mainly “local” context (meaning information from inputs $x^{(t')}$ where t' is not too far from t).

In the next part, you will build a more complex LSTM model, which is better at addressing vanishing gradients. The LSTM will be better able to remember a piece of information and keep it saved for many timesteps.

1.4 2 - Long Short-Term Memory (LSTM) network

This following figure shows the operations of an LSTM-cell.

Figure 4: LSTM-cell. This tracks and updates a “cell state” or memory variable $c^{(t)}$ at every time-step, which can be different from $a^{(t)}$.

Similar to the RNN example above, you will start by implementing the LSTM cell for a single time-step. Then you can iteratively call it from inside a for-loop to have it process an input with T_x time-steps.

1.4.1 About the gates

- Forget gate For the sake of this illustration, let’s assume we are reading words in a piece of text, and want use an LSTM to keep track of grammatical structures, such as whether the subject is singular or plural. If the subject changes from a singular word to a plural word, we need to find a way to get rid of our previously stored memory value of the singular/plural state. In an LSTM, the forget gate let’s us do this:

$$\Gamma_f^{(t)} = \sigma(W_f[a^{(t-1)}, x^{(t)}] + b_f) \quad (1)$$

Here, W_f are weights that govern the forget gate’s behavior. We concatenate $[a^{(t-1)}, x^{(t)}]$ and multiply by W_f . The equation above results in a vector $\Gamma_f^{(t)}$ with values between 0 and 1. This forget gate vector will be multiplied element-wise by the previous cell state $c^{(t-1)}$. So if one of the values of $\Gamma_f^{(t)}$ is 0 (or close to 0) then it means that the LSTM should remove that piece of information (e.g. the singular subject) in the corresponding component of $c^{(t-1)}$. If one of the values is 1, then it will keep the information.

- Update gate Once we forget that the subject being discussed is singular, we need to find a way to update it to reflect that the new subject is now plural. Here is the formula for the update gate:

$$\Gamma_u^{(t)} = \sigma(W_u[a^{(t-1)}, x^{(t)}] + b_u) \quad (2)$$

Similar to the forget gate, here $\Gamma_u^{(t)}$ is again a vector of values between 0 and 1. This will be multiplied element-wise with $\tilde{c}^{(t)}$, in order to compute $c^{(t)}$.

- Updating the cell To update the new subject we need to create a new vector of numbers that we can add to our previous cell state. The equation we use is:

$$\tilde{c}^{(t)} = \tanh(W_c[a^{(t-1)}, x^{(t)}] + b_c) \quad (3)$$

Finally, the new cell state is:

$$c^{(t)} = \Gamma_f^{(t)} * c^{(t-1)} + \Gamma_u^{(t)} * \tilde{c}^{(t)} \quad (4)$$

- Output gate To decide which outputs we will use, we will use the following two formulas:

$$\Gamma_o^{(t)} = \sigma(W_o[a^{(t-1)}, x^{(t)}] + b_o) \quad (5)$$

$$a^{(t)} = \Gamma_o^{(t)} * \tanh(c^{(t)}) \quad (6)$$

Where in equation 5 you decide what to output using a sigmoid function and in equation 6 you multiply that by the tanh of the previous state.

1.4.2 2.1 - LSTM cell

Exercise: Implement the LSTM cell described in the Figure (4).

Instructions: 1. Concatenate $a^{(t-1)}$ and $x^{(t)}$ in a single matrix: $concat = \begin{bmatrix} a^{(t-1)} \\ x^{(t)} \end{bmatrix}$ 2. Compute all the formulas 1-6. You can use `sigmoid()` (provided) and `np.tanh()`. 3. Compute the prediction $y^{(t)}$. You can use `softmax()` (provided).

In [20]: # GRADED FUNCTION: lstm_cell_forward

```
def lstm_cell_forward(xt, a_prev, c_prev, parameters):
    """
    Implement a single forward step of the LSTM-cell as described in Figure (4)

    Arguments:
    xt -- your input data at timestep "t", numpy array of shape (n_x, m)
    a_prev -- Hidden state at timestep "t-1", numpy array of shape (n_a, m)
    c_prev -- Memory state at timestep "t-1", numpy array of shape (n_a, m)
    parameters -- python dictionary containing:
        Wf -- Weight matrix of the forget gate, numpy array of shape (n_a, n_x + n_h)
        bf -- Bias of the forget gate, numpy array of shape (n_a, 1)
        Wi -- Weight matrix of the update gate, numpy array of shape (n_a, n_x + n_h)
        bi -- Bias of the update gate, numpy array of shape (n_a, 1)
        Wc -- Weight matrix of the first "tanh", numpy array of shape (n_a, n_x + n_h)
        bc -- Bias of the first "tanh", numpy array of shape (n_a, 1)
    """
```

```

        Wo -- Weight matrix of the output gate, numpy array of shape (n_a, m)
        bo -- Bias of the output gate, numpy array of shape (n_a, 1)
        Wy -- Weight matrix relating the hidden-state to the output, numpy array of shape (n_a, m)
        by -- Bias relating the hidden-state to the output, numpy array of shape (n_a, 1)

Returns:
a_next -- next hidden state, of shape (n_a, m)
c_next -- next memory state, of shape (n_a, m)
yt_pred -- prediction at timestep "t", numpy array of shape (n_y, m)
cache -- tuple of values needed for the backward pass, contains (a_next, c_next, a_prev, c_prev)

Note: ft/it/ot stand for the forget/update/output gates, cct stands for the candidate cell state,
      c stands for the memory value
"""

# Retrieve parameters from "parameters"
Wf = parameters["Wf"]
bf = parameters["bf"]
Wi = parameters["Wi"]
bi = parameters["bi"]
Wc = parameters["Wc"]
bc = parameters["bc"]
Wo = parameters["Wo"]
bo = parameters["bo"]
Wy = parameters["Wy"]
by = parameters["by"]

# Retrieve dimensions from shapes of xt and Wy
n_x, m = xt.shape
n_y, n_a = Wy.shape

### START CODE HERE ###
# Concatenate a_prev and xt (3 lines)
concat = np.zeros(shape=(n_a+n_x, m))
concat[: n_a, :] = a_prev
concat[n_a :, :] = xt

# Compute values for ft, it, cct, c_next, ot, a_next using the formulas given figure 4.1
s = sigmoid
h = np.tanh
ft = s(np.dot(Wf, concat) + bf)
it = s(np.dot(Wi, concat) + bi)
cct = h(np.dot(Wc, concat) + bc)
c_next = it*cct + ft*c_prev
ot = s(np.dot(Wo, concat) + bo)
a_next = ot * h(c_next)

# Compute prediction of the LSTM cell (1 line)

```



```

yt_pred = softmax(np.dot(Wy, a_next)+by)
### END CODE HERE ###

# store values needed for backward propagation in cache
cache = (a_next, c_next, a_prev, c_prev, ft, it, cct, ot, xt, parameters)

return a_next, c_next, yt_pred, cache

```

```

In [21]: np.random.seed(1)
xt = np.random.randn(3,10)
a_prev = np.random.randn(5,10)
c_prev = np.random.randn(5,10)
Wf = np.random.randn(5, 5+3)
bf = np.random.randn(5,1)
Wi = np.random.randn(5, 5+3)
bi = np.random.randn(5,1)
Wo = np.random.randn(5, 5+3)
bo = np.random.randn(5,1)
Wc = np.random.randn(5, 5+3)
bc = np.random.randn(5,1)
Wy = np.random.randn(2,5)
by = np.random.randn(2,1)

parameters = {"Wf": Wf, "Wi": Wi, "Wo": Wo, "Wc": Wc, "Wy": Wy, "bf": bf, "bi": bi, "bo": bo, "bc": bc, "by": by}

a_next, c_next, yt, cache = lstm_cell_forward(xt, a_prev, c_prev, parameters)
print("a_next[4] = ", a_next[4])
print("a_next.shape = ", c_next.shape)
print("c_next[2] = ", c_next[2])
print("c_next.shape = ", c_next.shape)
print("yt[1] = ", yt[1])
print("yt.shape = ", yt.shape)
print("cache[1][3] = ", cache[1][3])
print("len(cache) = ", len(cache))

a_next[4] = [-0.66408471  0.0036921  0.02088357  0.22834167 -0.85575339  0.00138482
 0.76566531  0.34631421 -0.00215674  0.43827275]
a_next.shape = (5, 10)
c_next[2] = [ 0.63267805  1.00570849  0.35504474  0.20690913 -1.64566718  0.11832942
 0.76449811 -0.0981561  -0.74348425 -0.26810932]
c_next.shape = (5, 10)
yt[1] = [ 0.79913913  0.15986619  0.22412122  0.15606108  0.97057211  0.31146381
 0.00943007  0.12666353  0.39380172  0.07828381]
yt.shape = (2, 10)
cache[1][3] = [-0.16263996  1.03729328  0.72938082 -0.54101719  0.02752074 -0.30821874
 0.07651101 -1.03752894  1.41219977 -0.37647422]
len(cache) = 10

```

```

Expected Output:
a_next[4]:
[-0.66408471 0.0036921 0.02088357 0.22834167 -0.85575339 0.00138482 0.76566531 0.34631421 -
0.00215674 0.43827275]
a_next.shape:
(5, 10)
c_next[2]:
[ 0.63267805 1.00570849 0.35504474 0.20690913 -1.64566718 0.11832942 0.76449811 -0.0981561
-0.74348425 -0.26810932]
c_next.shape:
(5, 10)
yt[1]:
[ 0.79913913 0.15986619 0.22412122 0.15606108 0.97057211 0.31146381 0.00943007 0.12666353
0.39380172 0.07828381]
yt.shape:
(2, 10)
cache[1][3]:
[-0.16263996 1.03729328 0.72938082 -0.54101719 0.02752074 -0.30821874 0.07651101 -1.03752894
1.41219977 -0.37647422]
len(cache):
10

```

1.4.3 2.2 - Forward pass for LSTM

Now that you have implemented one step of an LSTM, you can now iterate this over this using a for-loop to process a sequence of T_x inputs.

Figure 5: LSTM over multiple time-steps.

Exercise: Implement `lstm_forward()` to run an LSTM over T_x time-steps.

Note: $c^{(0)}$ is initialized with zeros.

```
In [22]: # GRADED FUNCTION: lstm_forward
```

```

def lstm_forward(x, a0, parameters):
    """
    Implement the forward propagation of the recurrent neural network using an LSTM-cell.

    Arguments:
    x -- Input data for every time-step, of shape (n_x, m, T_x).
    a0 -- Initial hidden state, of shape (n_a, m)
    parameters -- python dictionary containing:
        Wf -- Weight matrix of the forget gate, numpy array of shape (n_a, n_x)
        bf -- Bias of the forget gate, numpy array of shape (n_a, 1)
        Wi -- Weight matrix of the update gate, numpy array of shape (n_a, n_x)
        bi -- Bias of the update gate, numpy array of shape (n_a, 1)
        Wc -- Weight matrix of the first "tanh", numpy array of shape (n_a, n_x)
        bc -- Bias of the first "tanh", numpy array of shape (n_a, 1)
        Wo -- Weight matrix of the output gate, numpy array of shape (n_a, n_x)
        bo -- Bias of the output gate, numpy array of shape (n_a, 1)
    """

```

Wy -- Weight matrix relating the hidden-state to the output, numpy array
by -- Bias relating the hidden-state to the output, numpy array

Returns:

a -- Hidden states for every time-step, numpy array of shape (n_a, m, T_x)
y -- Predictions for every time-step, numpy array of shape (n_y, m, T_x)
caches -- tuple of values needed for the backward pass, contains (list of all the caches, x)
 """

Initialize "caches", which will track the list of all the caches
 caches = []

START CODE HERE

Retrieve dimensions from shapes of x and parameters['Wy'] (2 lines)

n_x, m, T_x = x.shape
 n_y, n_a = parameters['Wy'].shape

initialize "a", "c" and "y" with zeros (3 lines)

a = np.zeros(shape=(n_a, m, T_x))
 c = np.zeros(shape=(n_a, m, T_x))
 y = np.zeros(shape=(n_y, m, T_x))

Initialize a_next and c_next (2 lines)

a_next = a0
 c_next = np.zeros_like(a0)

loop over all time-steps

```
for t in range(T_x):
    # Update next hidden state, next memory state, compute the prediction, get the
    a_next, c_next, yt, cache = lstm_cell_forward(x[:, :, t], a_next, c_next, parameters)
    # Save the value of the new "next" hidden state in a (1 line)
    a[:, :, t] = a_next
    # Save the value of the prediction in y (1 line)
    y[:, :, t] = yt
    # Save the value of the next cell state (1 line)
    c[:, :, t] = c_next
    # Append the cache into caches (1 line)
    caches.append(cache)
```

END CODE HERE

store values needed for backward propagation in cache
 caches = (caches, x)

return a, y, c, caches

```
In [23]: np.random.seed(1)
         x = np.random.randn(3,10,7)
```

```

a0 = np.random.randn(5,10)
Wf = np.random.randn(5, 5+3)
bf = np.random.randn(5,1)
Wi = np.random.randn(5, 5+3)
bi = np.random.randn(5,1)
Wo = np.random.randn(5, 5+3)
bo = np.random.randn(5,1)
Wc = np.random.randn(5, 5+3)
bc = np.random.randn(5,1)
Wy = np.random.randn(2,5)
by = np.random.randn(2,1)

parameters = {"Wf": Wf, "Wi": Wi, "Wo": Wo, "Wc": Wc, "Wy": Wy, "bf": bf, "bi": bi, "bo": bo, "bc": bc, "by": by}

a, y, c, caches = lstm_forward(x, a0, parameters)
print("a[4][3][6] = ", a[4][3][6])
print("a.shape = ", a.shape)
print("y[1][4][3] =", y[1][4][3])
print("y.shape = ", y.shape)
print("caches[1][1][1] =", caches[1][1][1])
print("c[1][2][1]", c[1][2][1])
print("len(caches) = ", len(caches))

a[4][3][6] = 0.172117767533
a.shape = (5, 10, 7)
y[1][4][3] = 0.95087346185
y.shape = (2, 10, 7)
caches[1][1][1] = [ 0.82797464  0.23009474  0.76201118 -0.22232814 -0.20075807  0.18656139
 0.41005165]
c[1][2][1] -0.855544916718
len(caches) = 2

```

Expected Output:

```

a[4][3][6] =
0.172117767533
a.shape =
(5, 10, 7)
y[1][4][3] =
0.95087346185
y.shape =
(2, 10, 7)
caches[1][1][1] =
[ 0.82797464 0.23009474 0.76201118 -0.22232814 -0.20075807 0.18656139 0.41005165]
c[1][2][1] =
-0.855544916718
len(caches) =
2

```